

Machine Learning for Beginners



Who are we?




Christine Koppelt



Sonja Scheungrab



What is Machine Learning?



Definition 1:

**»Field of study that gives computers the ability to learn without being explicitly programmed.«
(Arthur Samuel, 1959)**

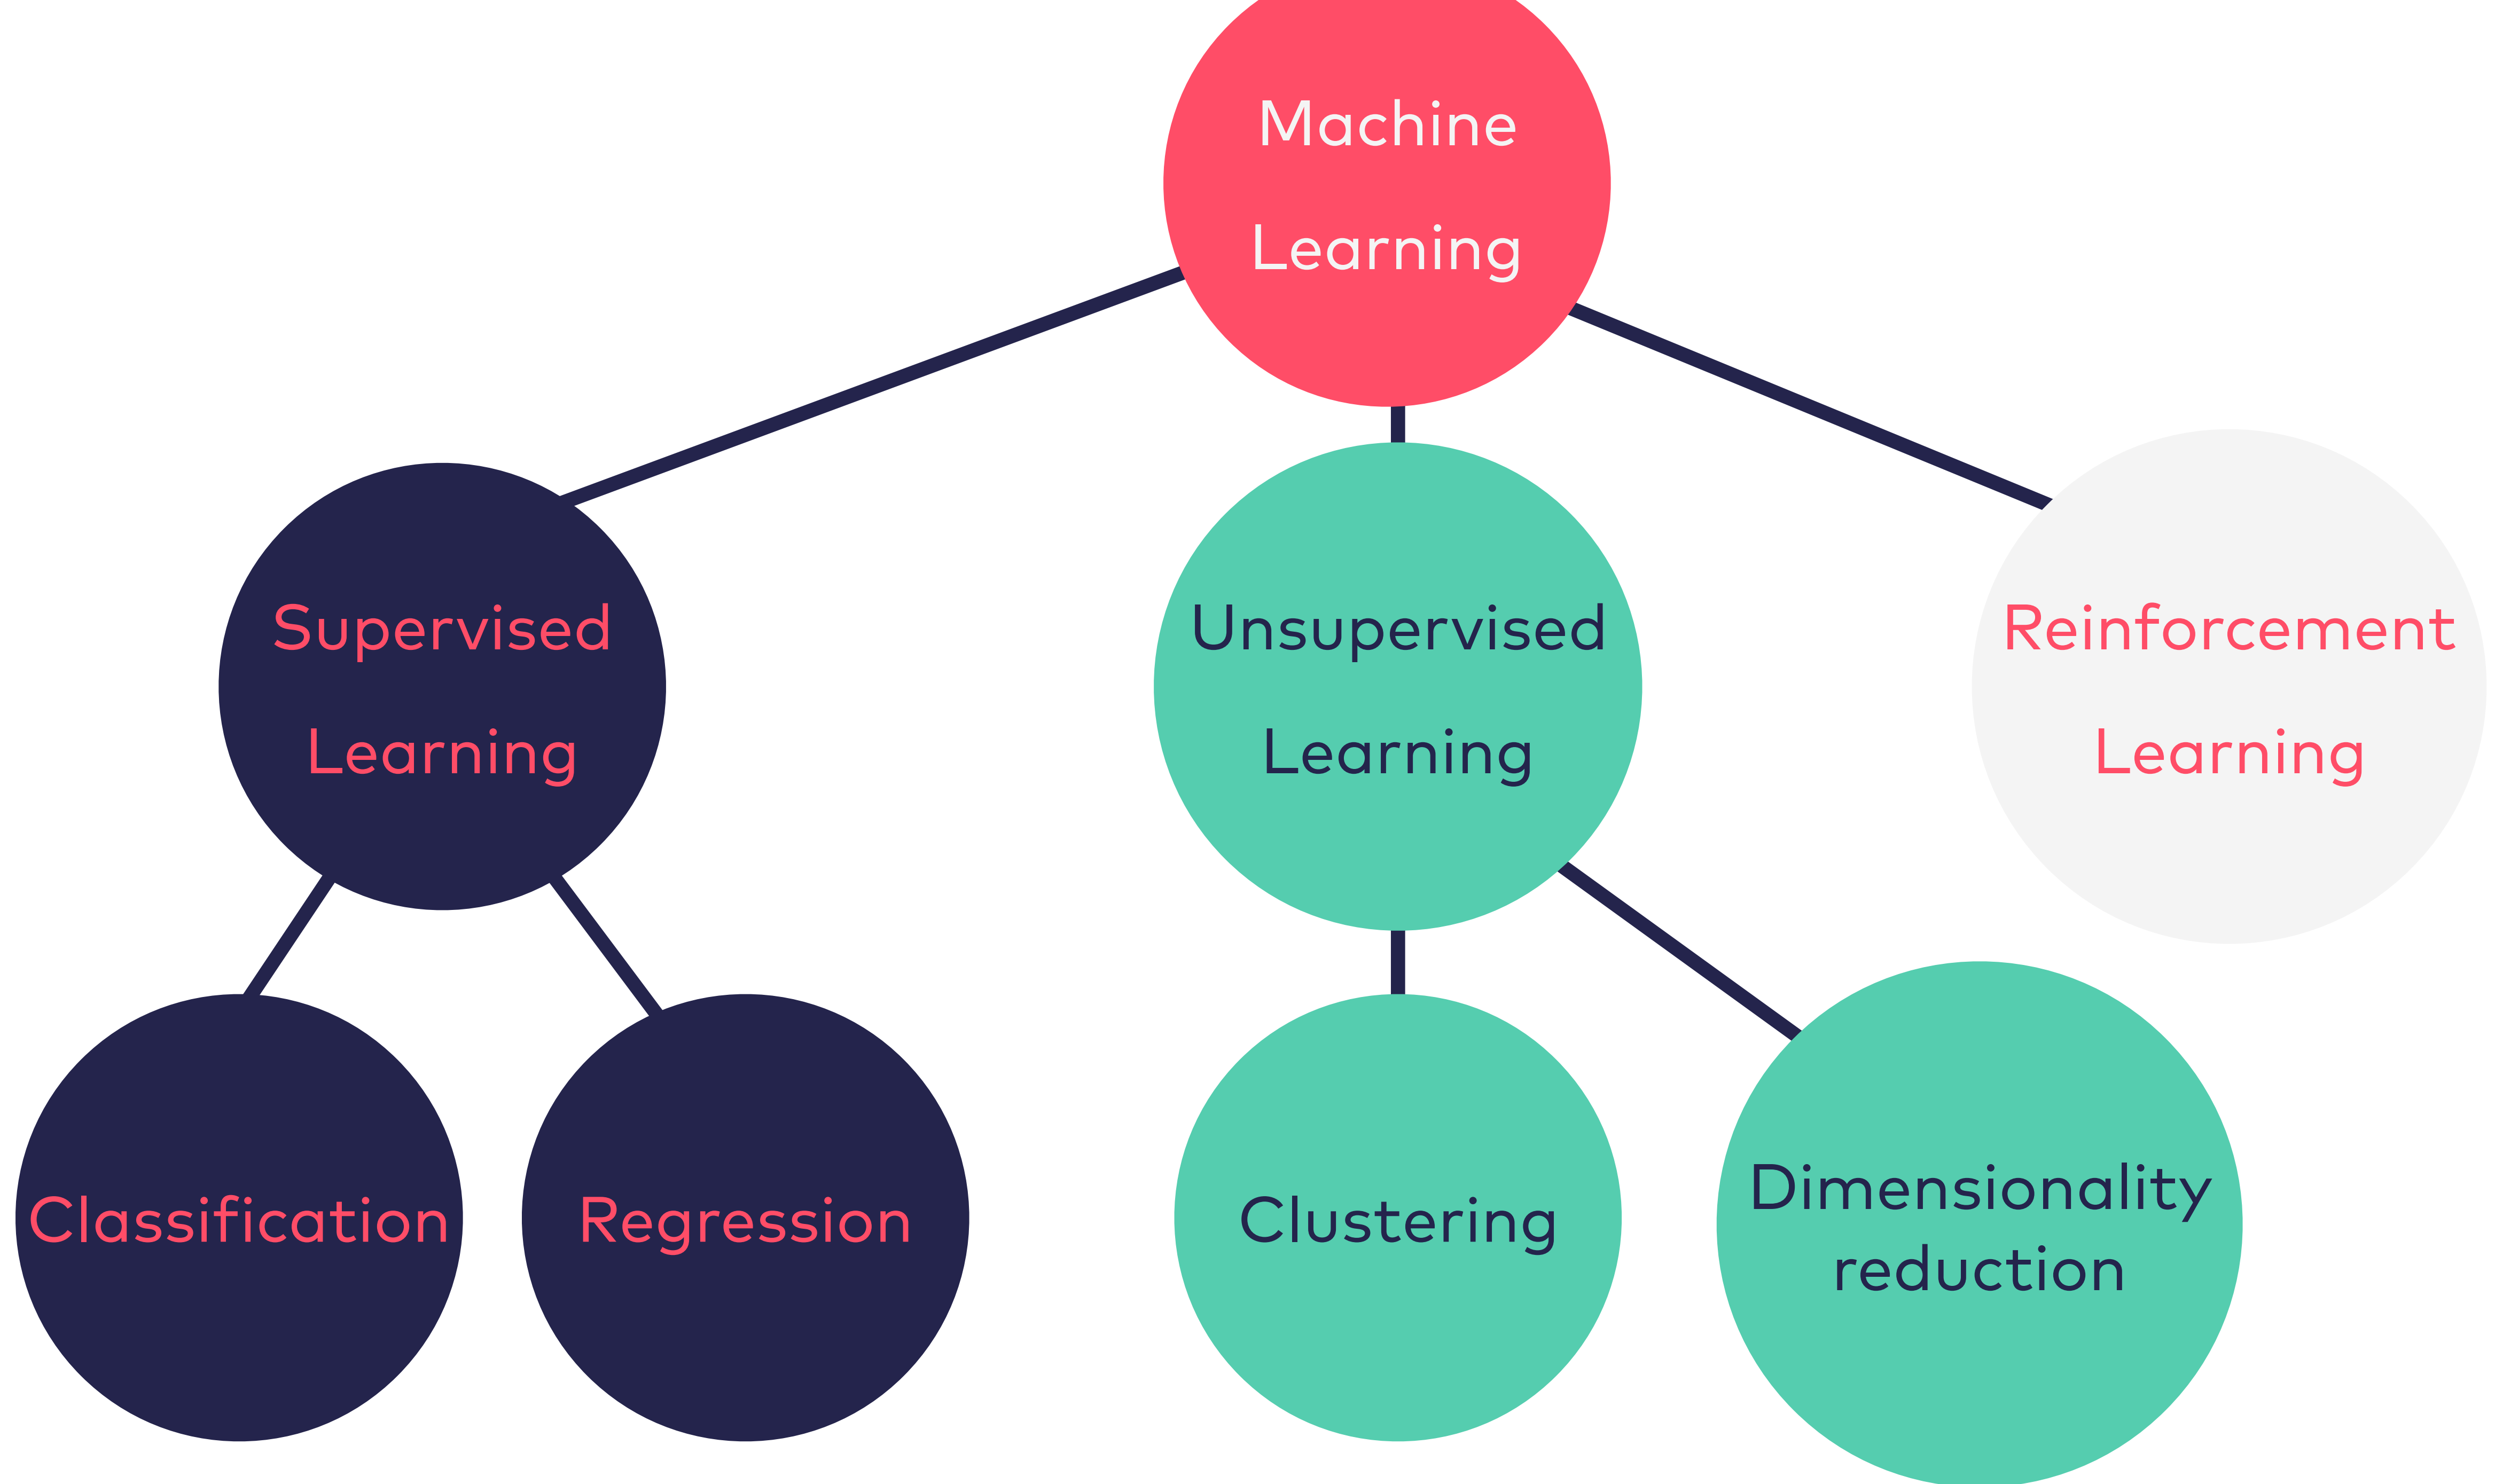
Definition 2:

»Well-posed Learning Problem: A computer program is said to learn from **experience E** with respect to some **task T** and some **performance measure P**, if its performance on **T**, as measured by **P**, improves with experience **E**.«

(Tom Mitchell, 1998)

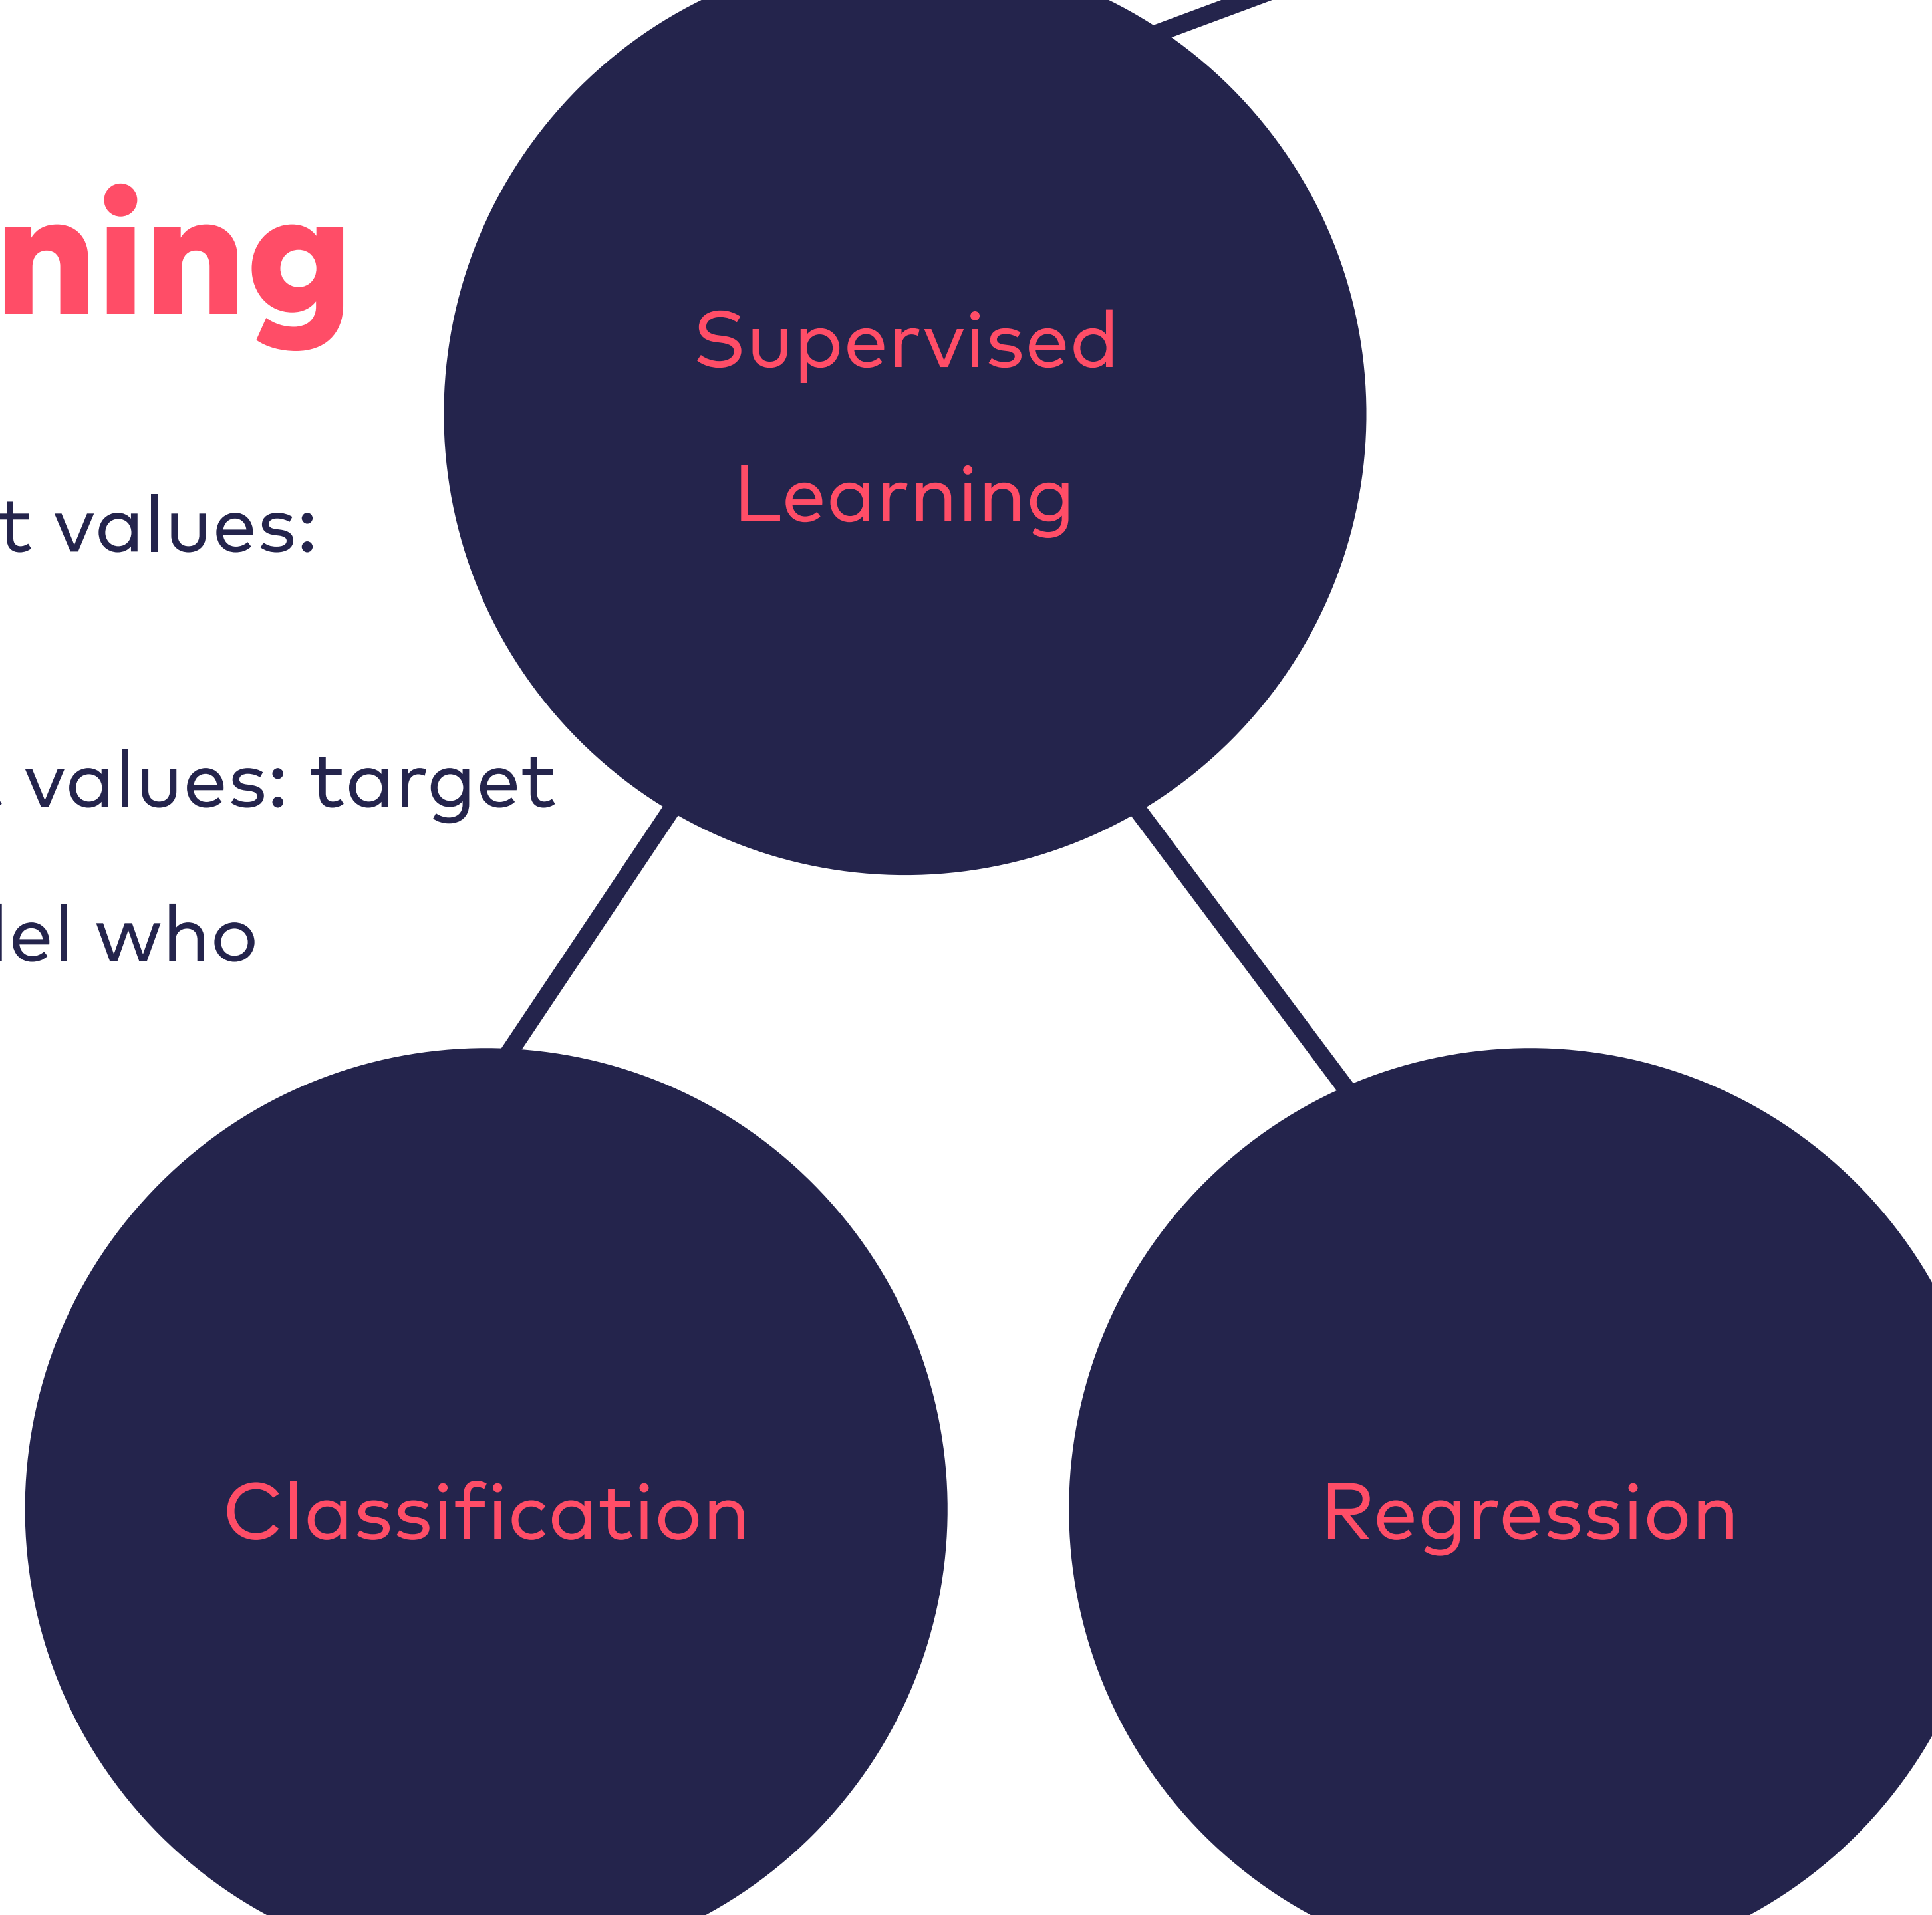
Plan for the Workshop

- Overview about the field of machine learning
- Focus on the practical work with useful frameworks
- Intuition about the background
- Analyze data together
- Experiment with different classification algorithms on that data.



Supervised Learning

- Dataset with input and output values: "Correct answer" is given
- Input values: Features, output values: target
- We are "teachers" of the model who point out errors



Unsupervised Learning

- No target values
- Find structure in dataset
- e.g clustering the data

Unsupervised Learning

Clustering

Dimensionality reduction

Reinforcement Learning

- We don't tell our model **what** it did wrong, only **if** it did good or bad
- System of reward (punishment = negative reward)
- With that guidance the model learns "by itself" trying to maximize the reward
- Takes up "natural" learning strategies (e.g. learning to walk)



Reinforcement
Learning

Classification

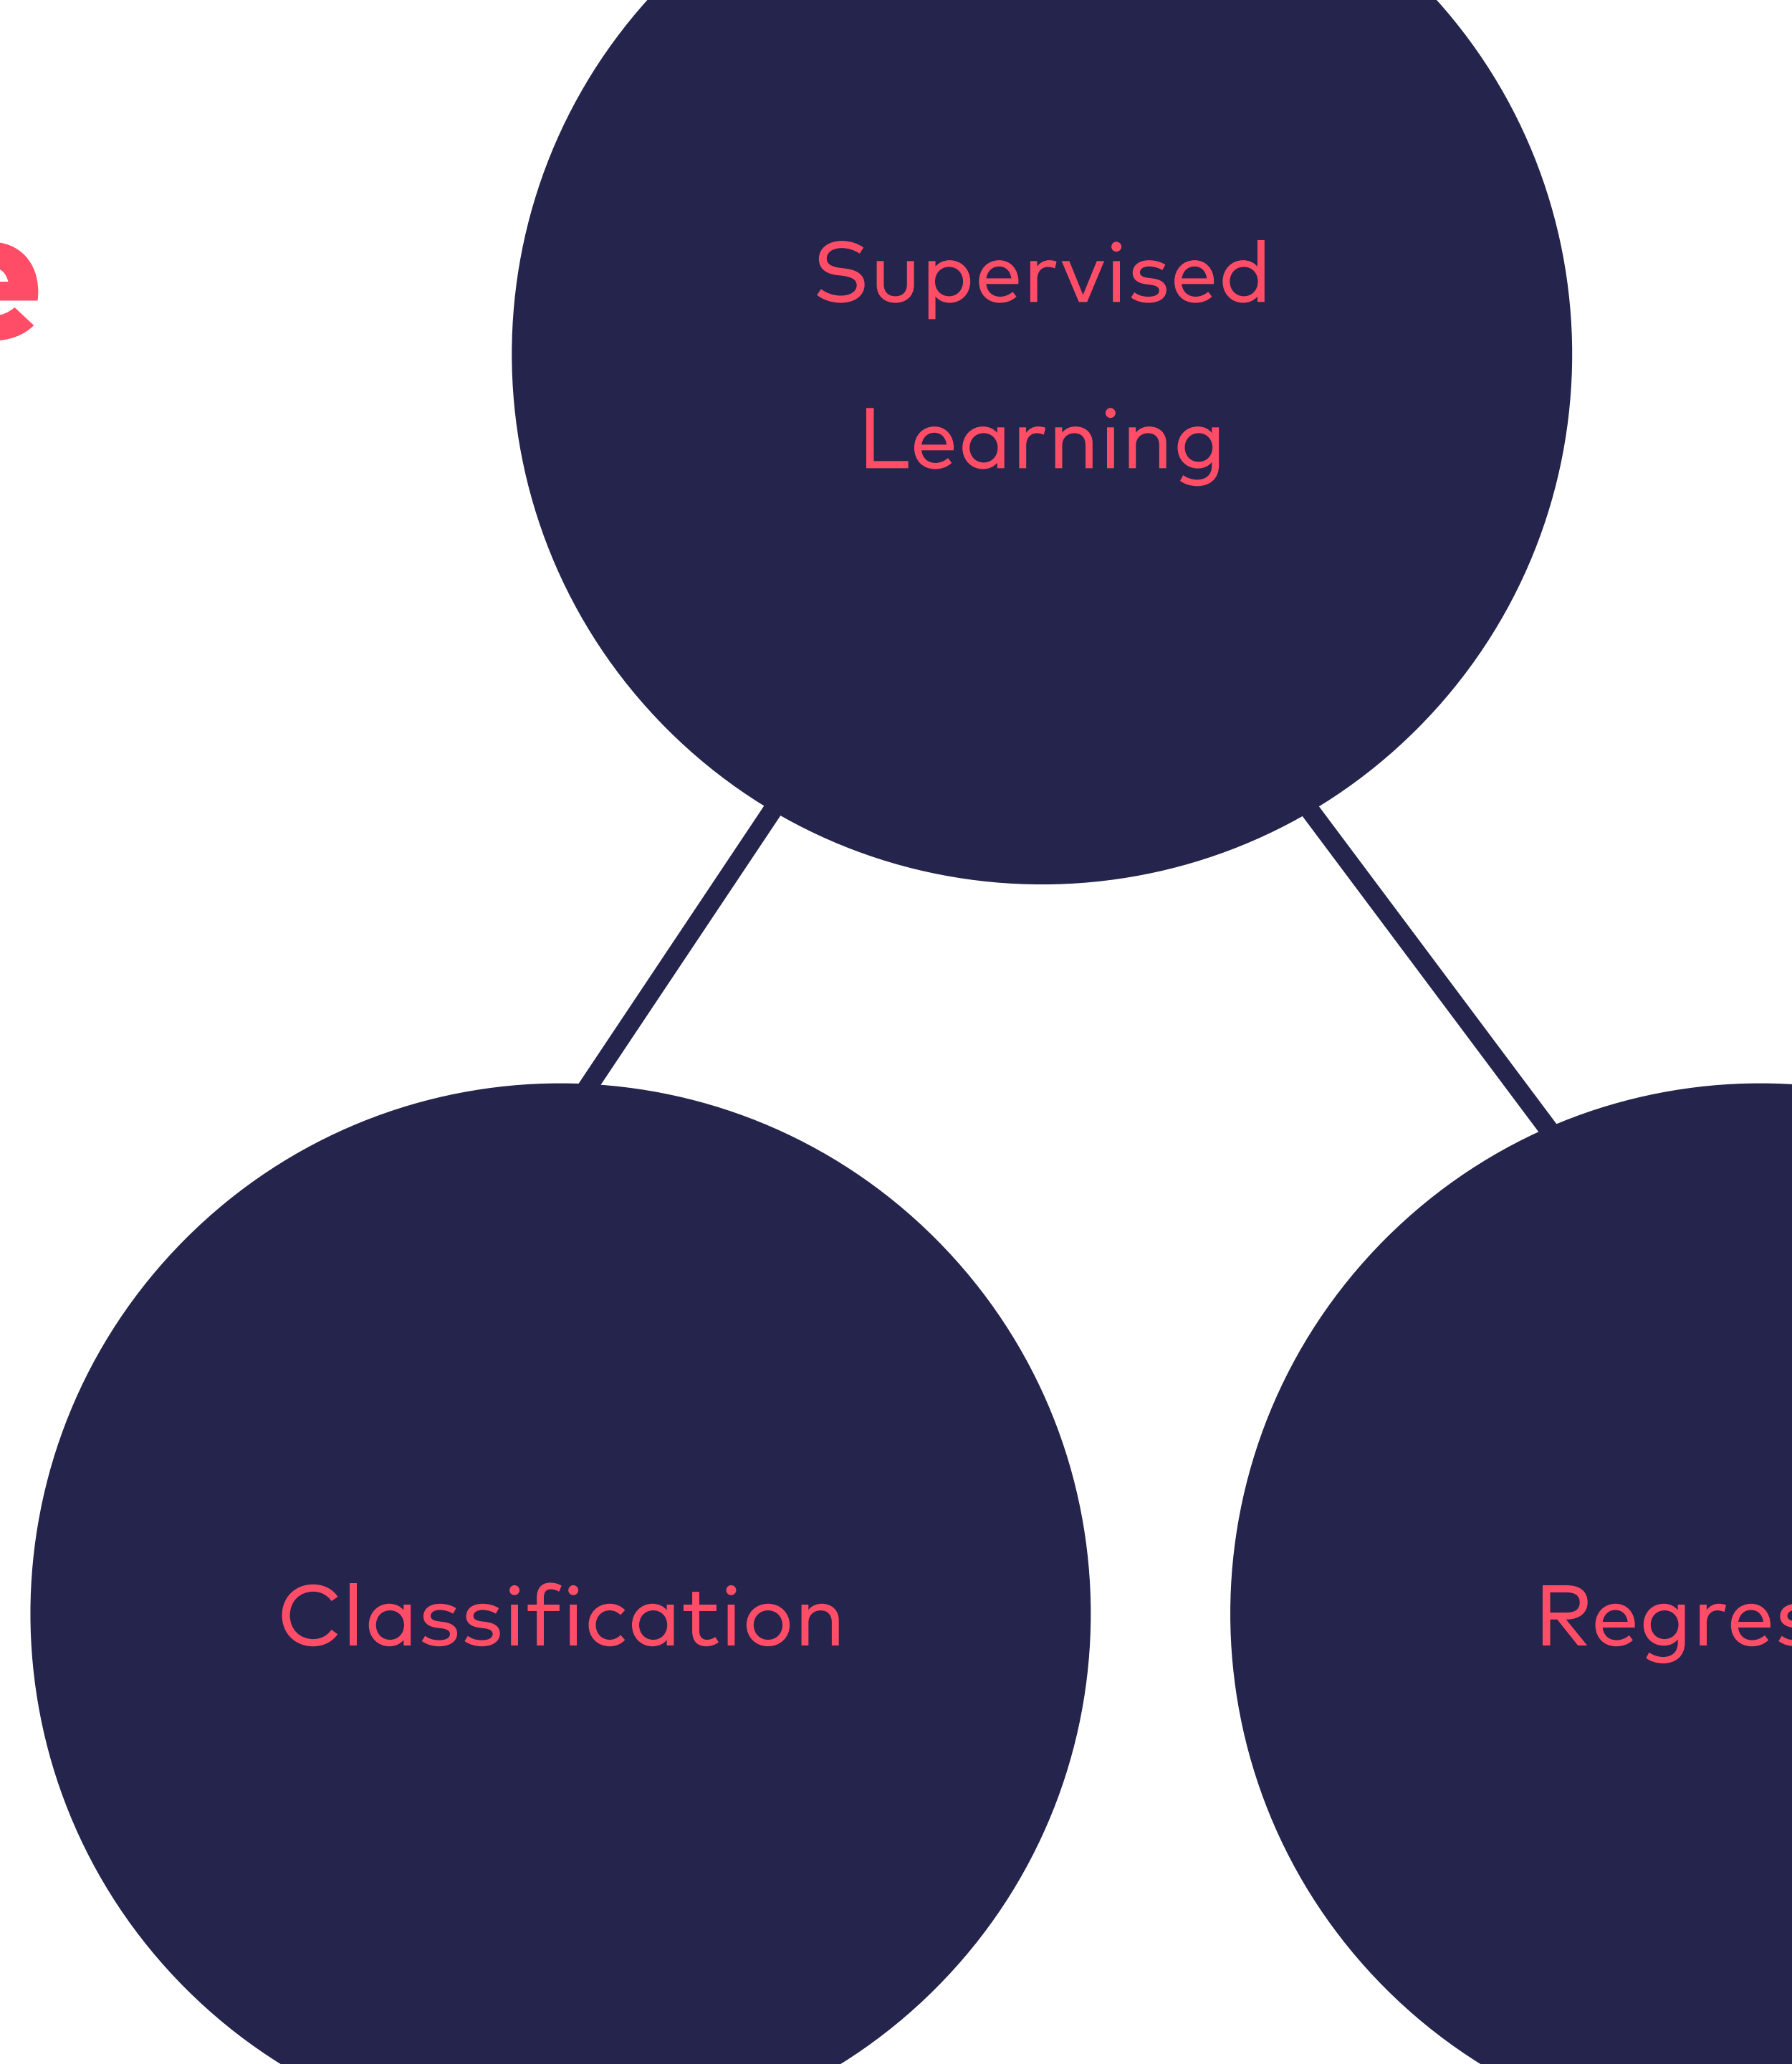
- Predict results in a discrete output
- Map input variables into two or more discrete categories: 🐱🐶🐭

Regression

- Predict results within a continuous output
- Map input variables to some continuous function
- House prizes, stock market...

The usual procedure

1. Choose your model
2. Instantiate your model with hyperparameters
3. Place your data in pattern matrix and a target vector
4. Train the model on the data = **fit**
5. Check the trained model on unknown (=test) data = **predict**



Learning Process

- Examine why errors happened during training
- Adapt the parameters/weights in order to get a better result
- Goal: Find a good general model we can use on unknown data

size = 70 cm, pattern = stripes >> 🐶?



pattern = stripes >> 🐯

size = 70 cm, pattern = stripes >> 🐯



What happens exactly?

Maths magic ✨

- Goal: find a "good" hypothesis function **h**
- The accuracy of our hypothesis is measured with a **cost function** like the **squared error function** that takes an average difference of all the results of the hypothesis with inputs from x 's (= y "hat") and the actual output y 's

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2 = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x_i) - y_i)^2$$

- We minimize the error with **gradient descent**: We take the derivative of our cost function, which tells us how to update our theta-weights

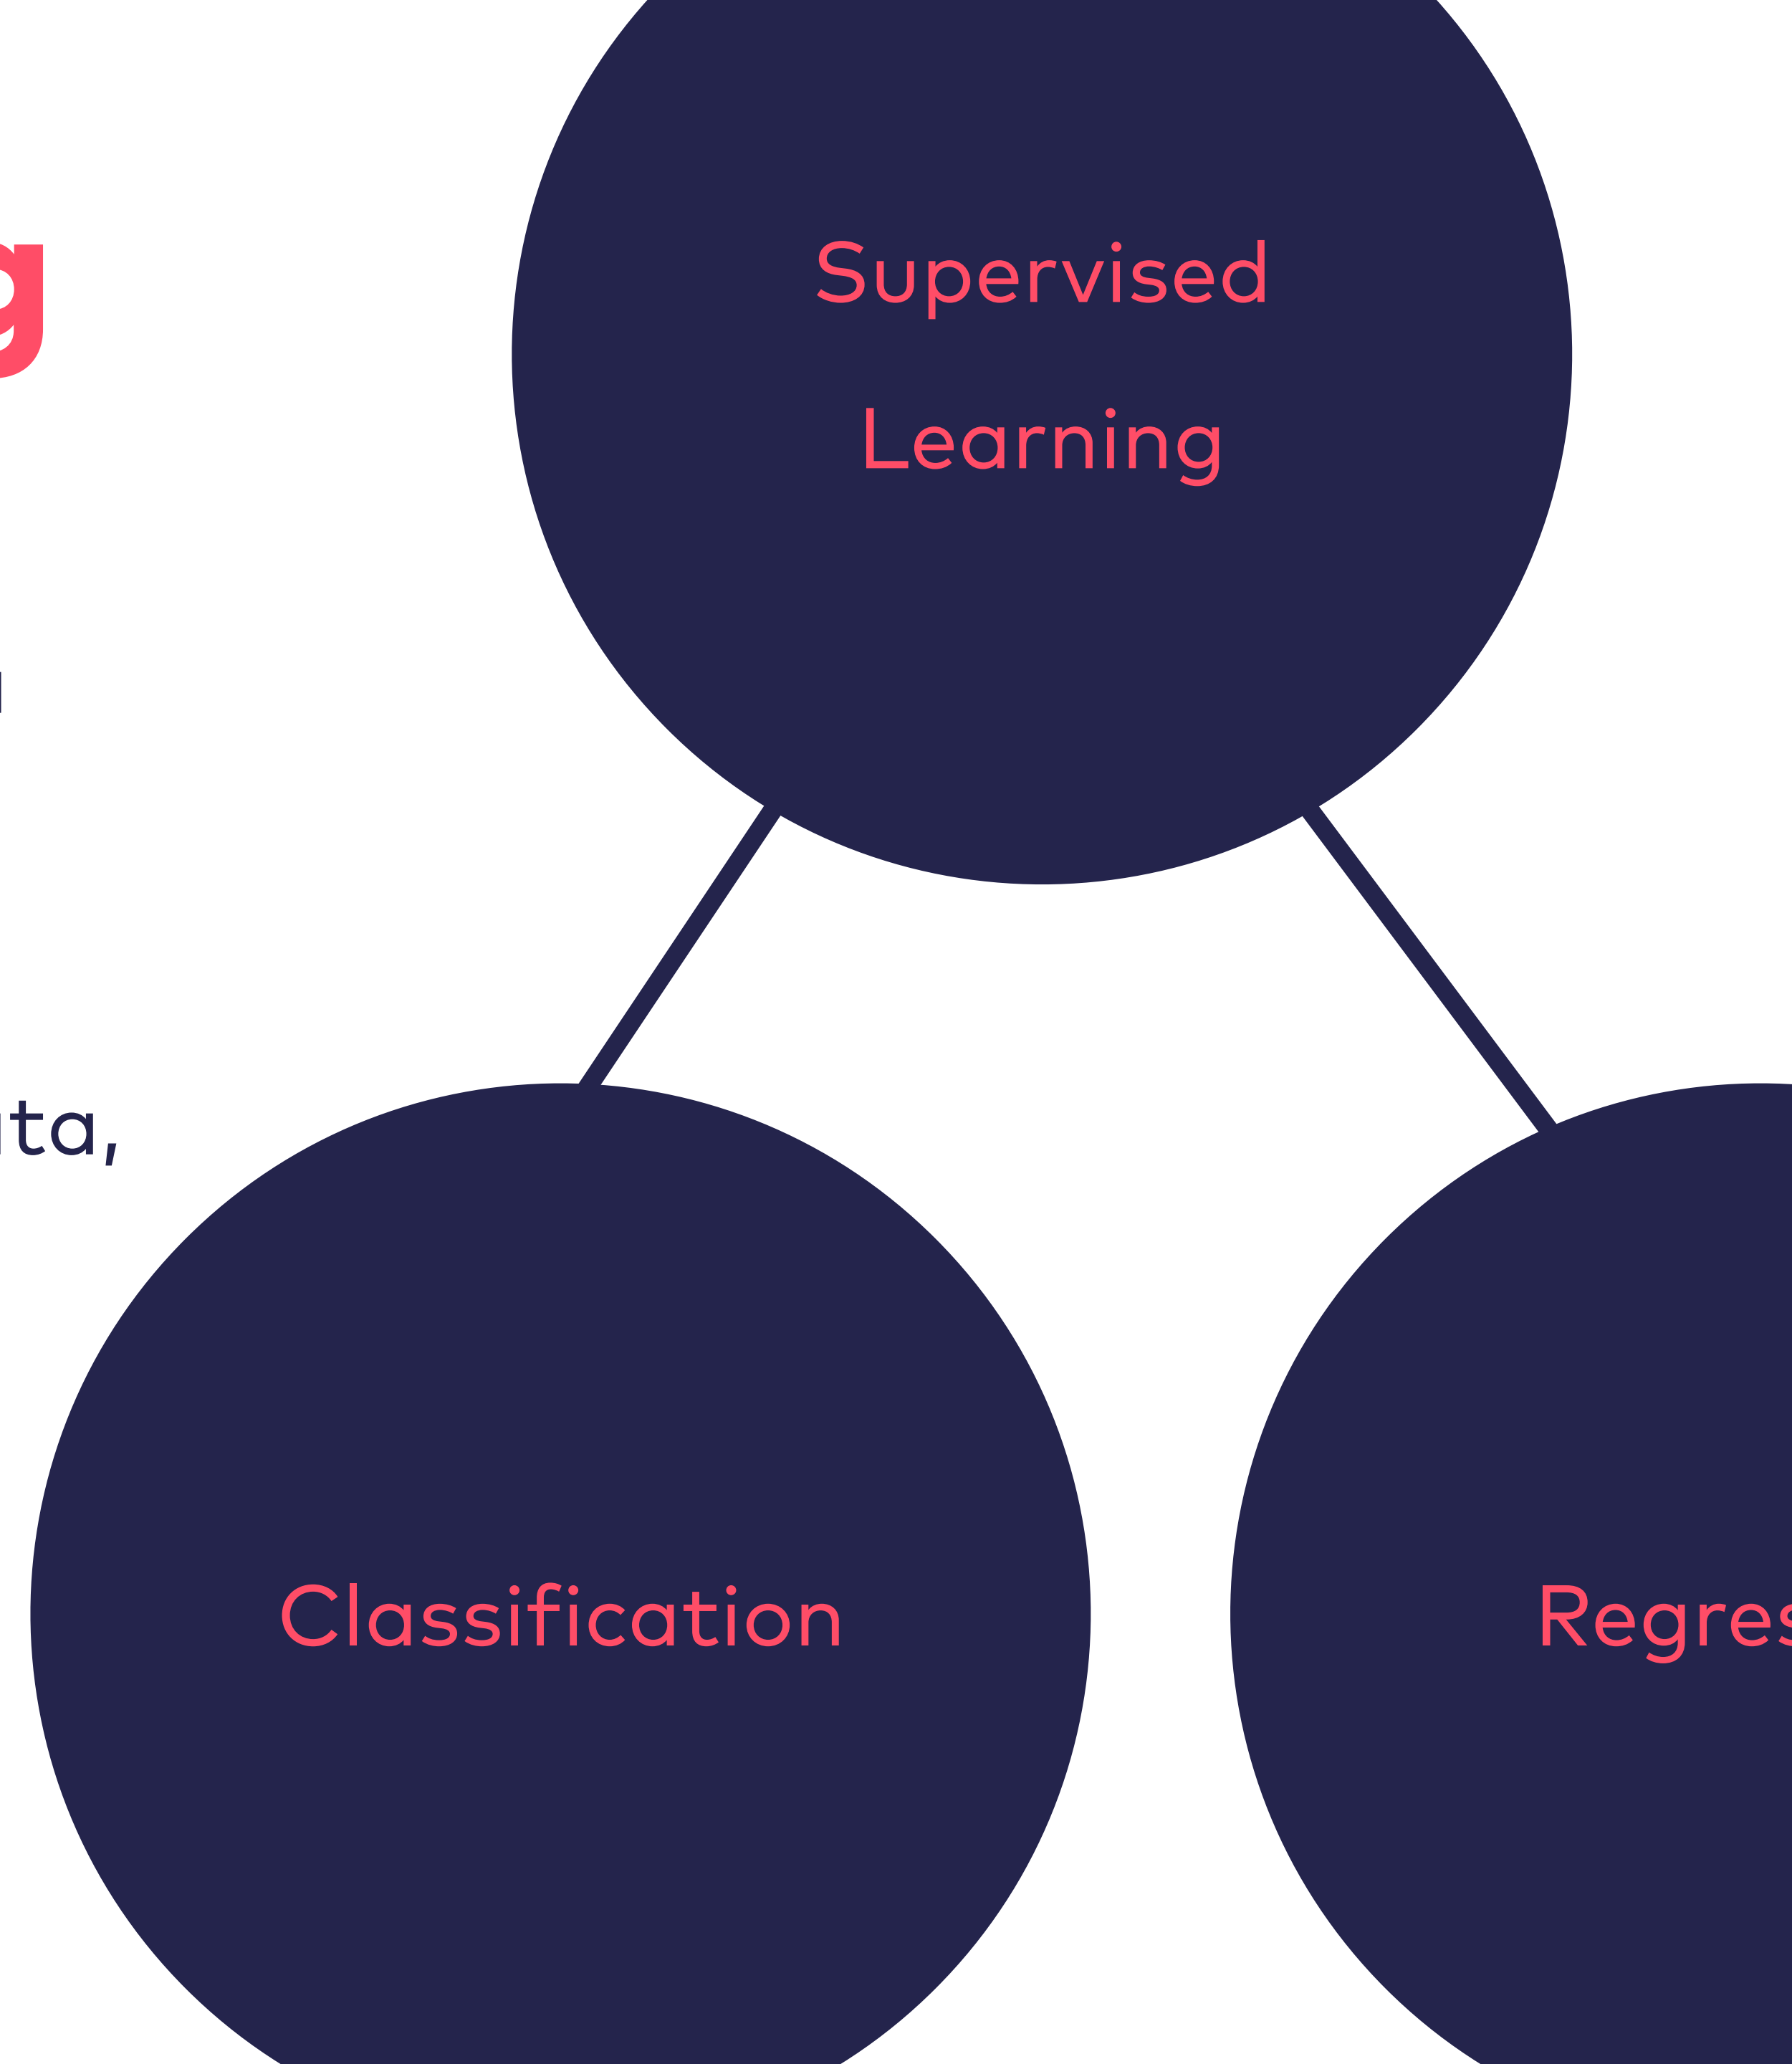


In other words

- Predict ...
- Check, how bad we did ...
- Minimize our error ...
- Repeat.

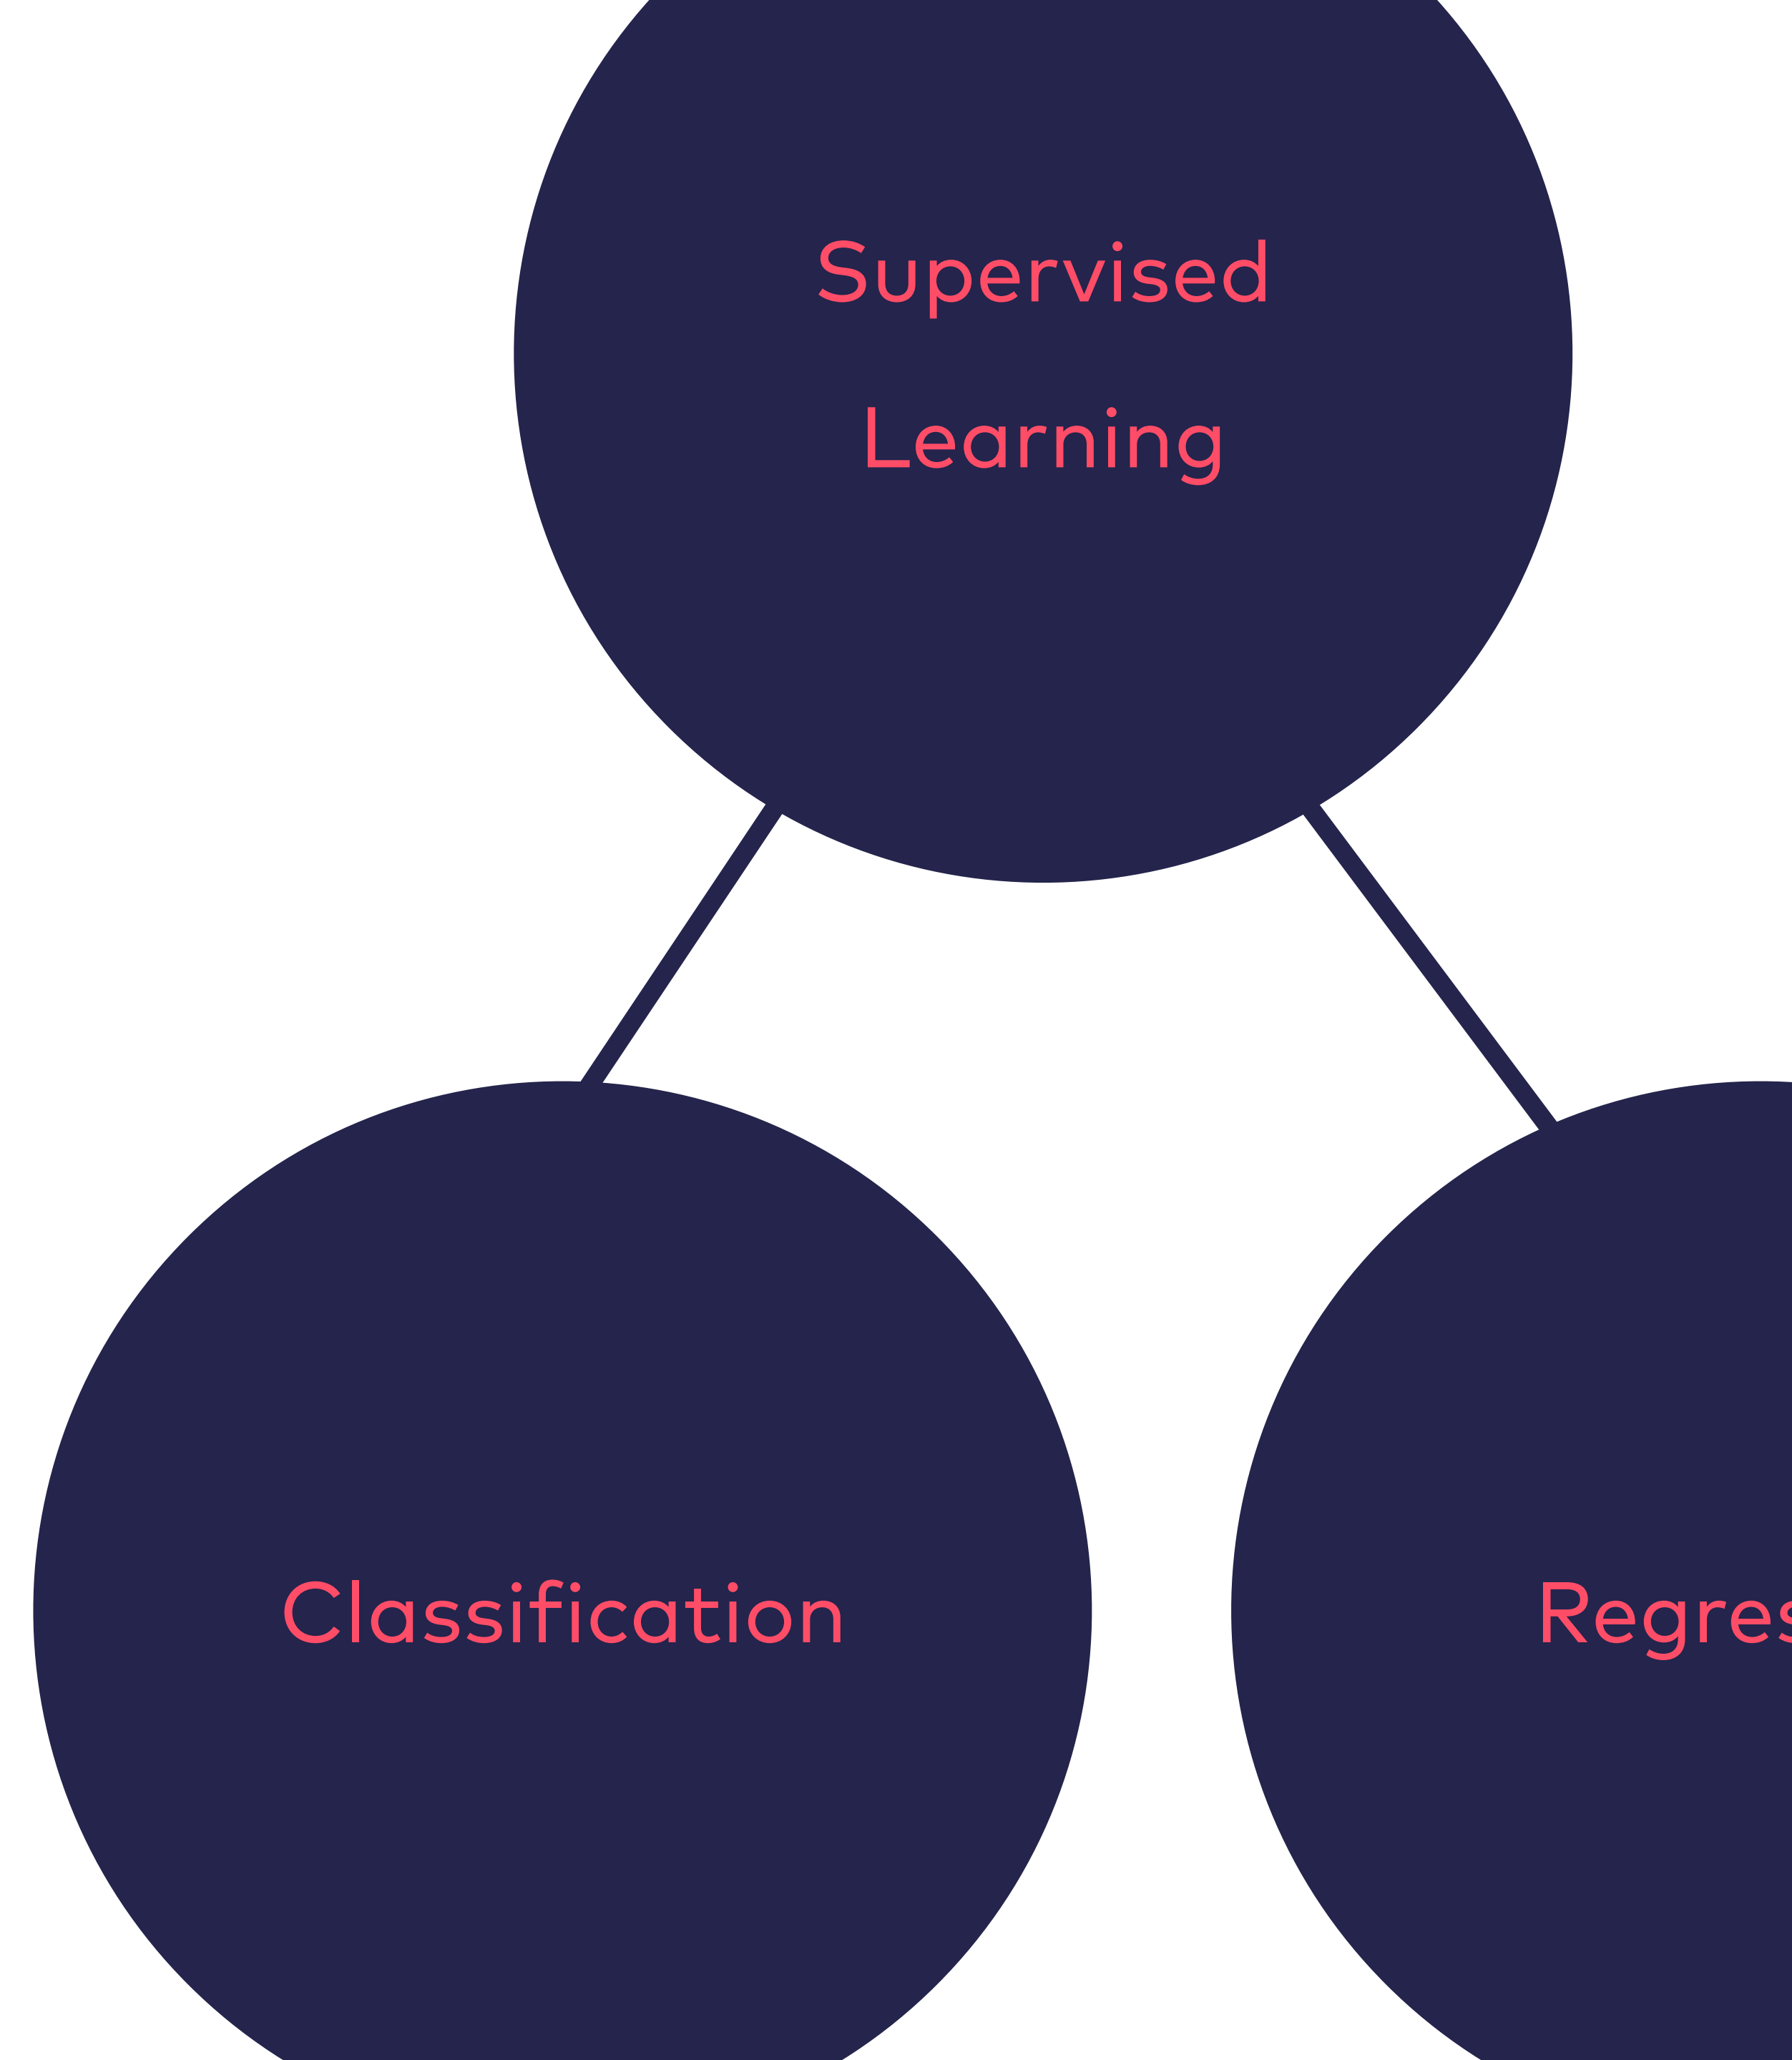
Training and Testing

- Split dataset in training and test sets
- e.g. 80 % training data, 20 % test data
- Remove the output/target values of test data
- "Train" the model with the trainings data, until we have a satisfying algorithm (good match with our data targets)
- Test the trained model on the unknown test data






Problems

- Overfitting
- Underfitting
- Bad data quality
- Small data quantity



Underfitting and Overfitting



Classification Algorithms

- Logistic Regression
- K-nearest neighbor
- Support Vector Machine
- Decision Tree / Random Forest

